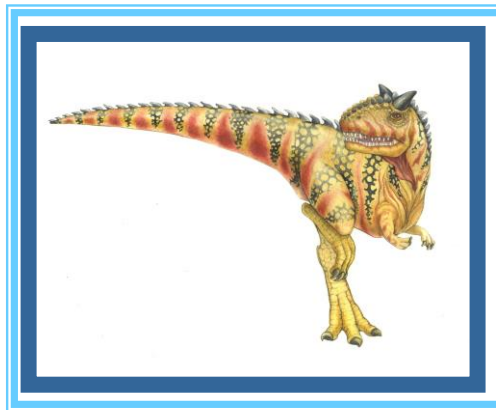


Chapter 8:

Memory-Management

Strategies



Chapter 8: Memory Management Strategies

Background

Swapping

Contiguous Memory Allocation

Paging

Structure of the Page Table

Segmentation

Example: The Intel Pentium

Objectives

To provide a detailed description of various ways of **organizing memory hardware**

To discuss various memory-management techniques, including **paging** and **segmentation**

To provide a detailed description of the Intel Pentium, which supports both **pure segmentation** and **segmentation with paging**

Background

Program must be brought (from disk) into **memory** and placed within a process for it to be run

Main memory and registers are only storage CPU can access directly

Register access in one CPU clock (or less)

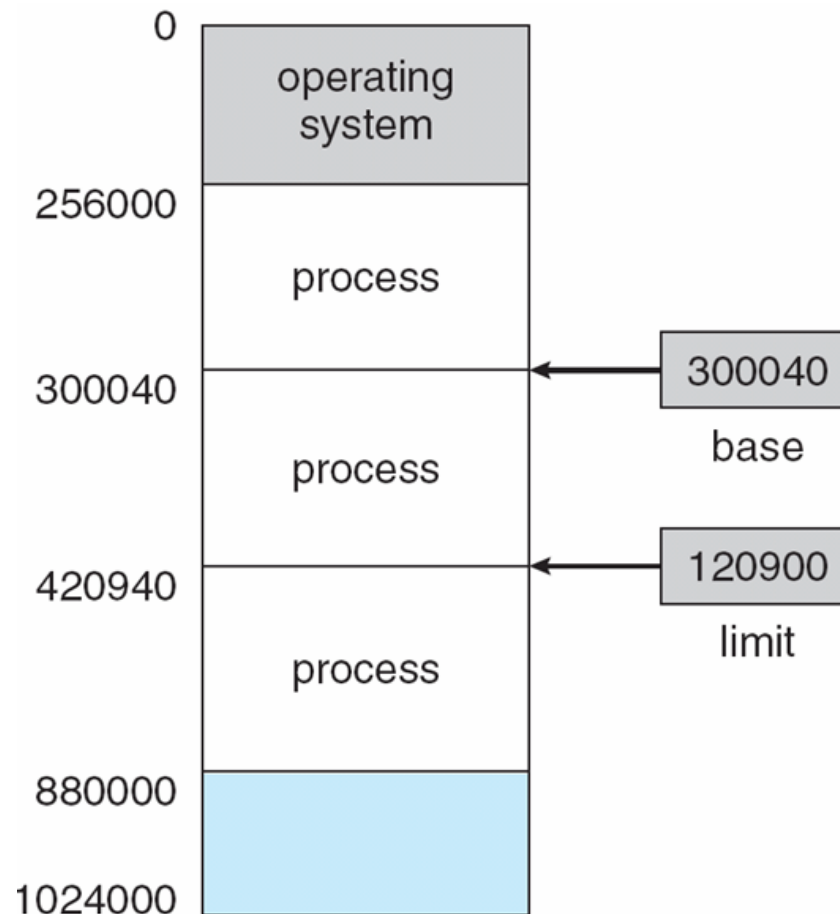
Main memory can take many cycles

Cache sits between main memory and CPU registers

Protection of memory is required to ensure correct operation

Base and Limit Registers

A pair of **base** and **limit** registers define the **logical address space**



Binding of Instructions and Data to Memory

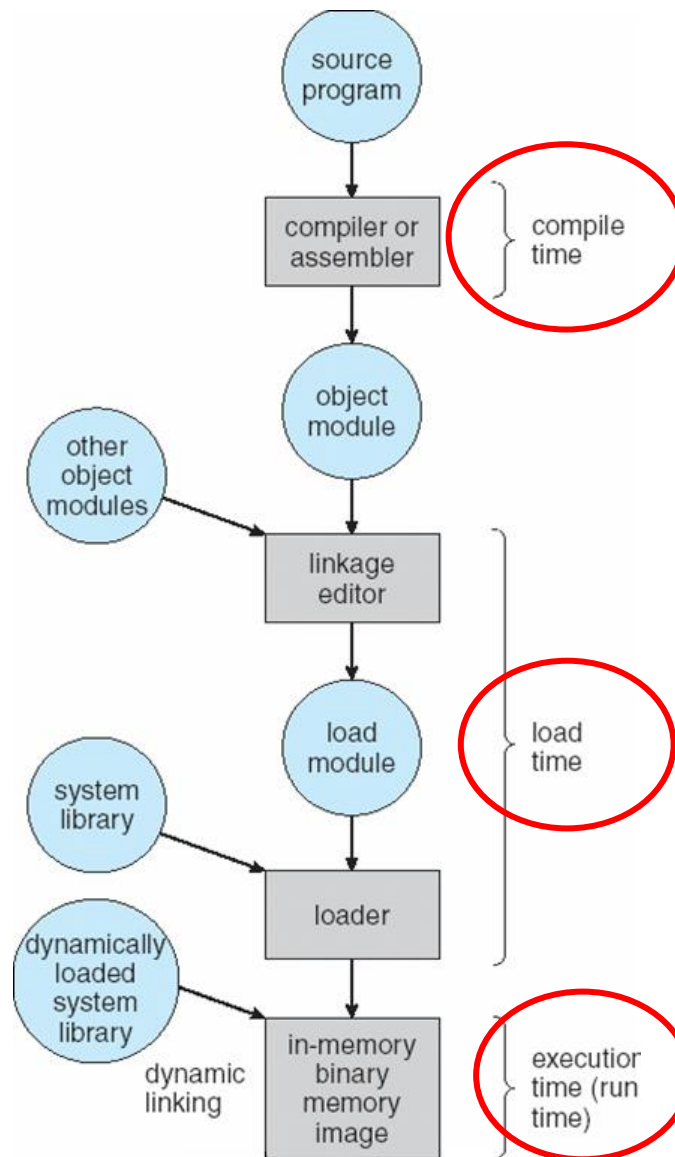
Address binding of instructions and data to memory addresses can happen at three different stages

Compile time: If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes

Load time: Must generate **relocatable code** if memory location is not known at compile time

Execution time: Binding delayed until run time if the process can be moved during its execution from one memory segment to another. Need hardware support for address maps (e.g., base and limit registers)

Multistep Processing of a User Program



Logical vs. Physical Address Space

The concept of a **logical address space** that is bound to a separate **physical address space** is central to proper memory management

Logical address – generated by the CPU; also referred to as **virtual address**

Physical address – address seen by the memory unit

Logical and physical addresses **are the same** in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses **differ in execution-time address-binding scheme**

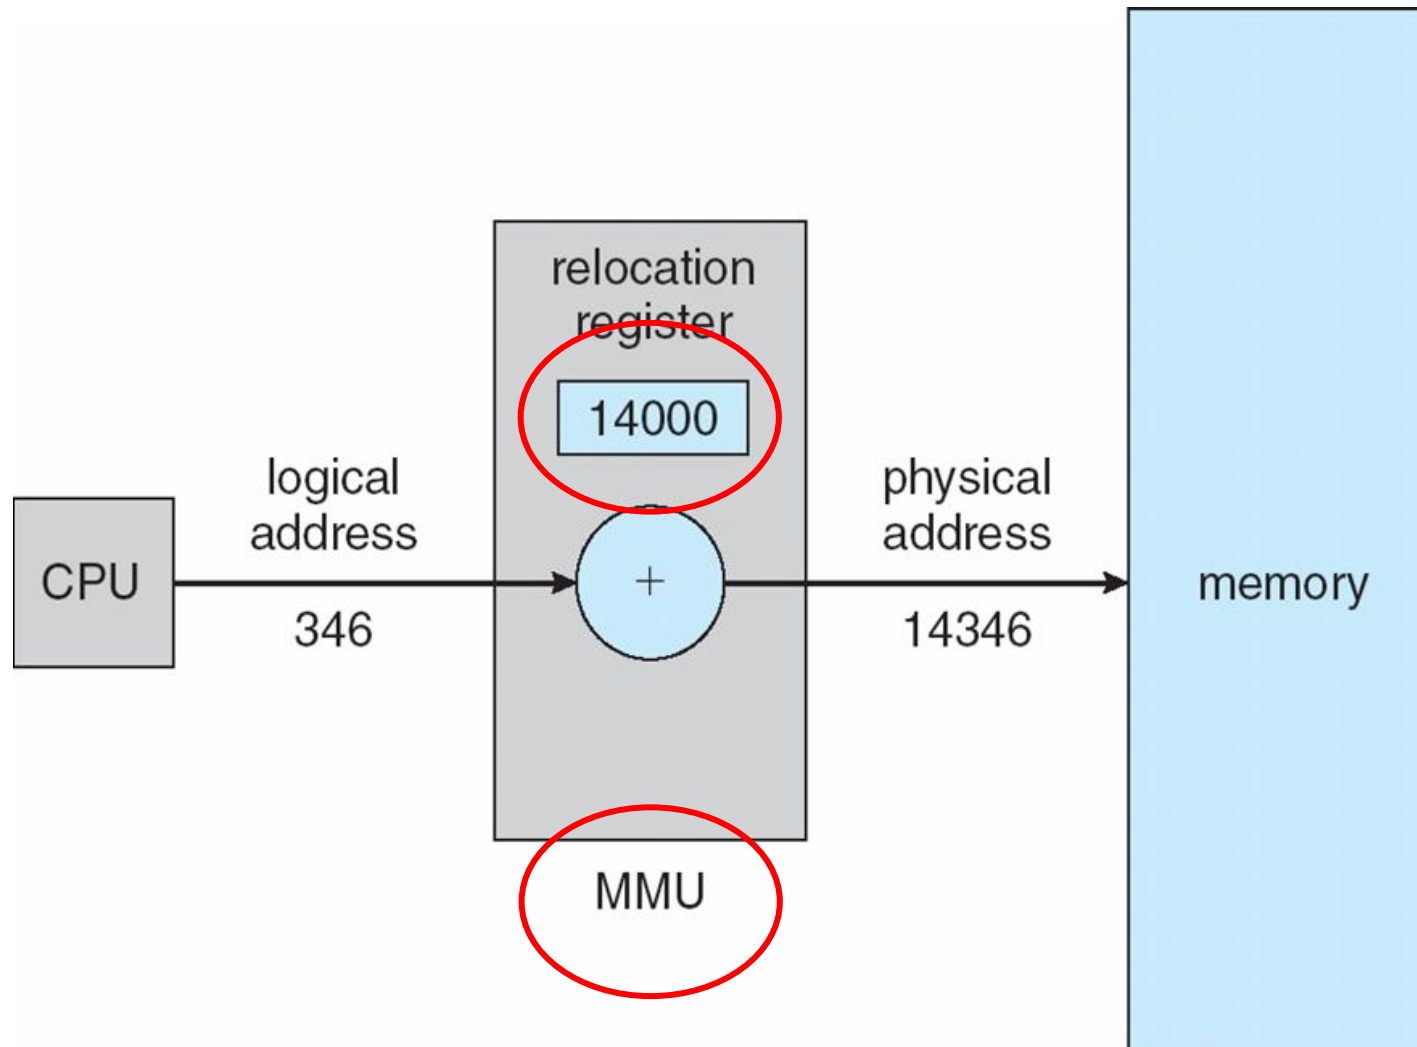
Memory-Management Unit (MMU)

Hardware device that **maps** virtual to physical address

In MMU scheme, the value in the **relocation register** is added to every address generated by a user process at the time it is sent to memory

The user program deals with **logical** addresses; it never sees the **real** physical addresses

Dynamic relocation using a relocation register



Dynamic Loading

Routine is not loaded until it is called

Better memory-space utilization; **unused routine is never loaded**

Useful when large amounts of code are needed to handle **infrequently occurring cases**

No special support from the operating system is required

Dynamic Linking

Linking postponed until execution time

Small piece of code, **stub**, used to locate the appropriate **memory-resident library routine**

Stub replaces itself with the address of the routine, and executes the routine

Operating system needed to check if routine is in processes' memory address

Dynamic linking is particularly **useful for libraries**

System also known as **shared libraries**

Swapping

A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution

Backing store – **fast disk large enough** to accommodate copies of all memory images for all users; must provide direct access to these memory images

Roll out, roll in – swapping variant used for **priority-based scheduling algorithms**;

lower-priority process is swapped out so higher-priority process can be loaded and executed

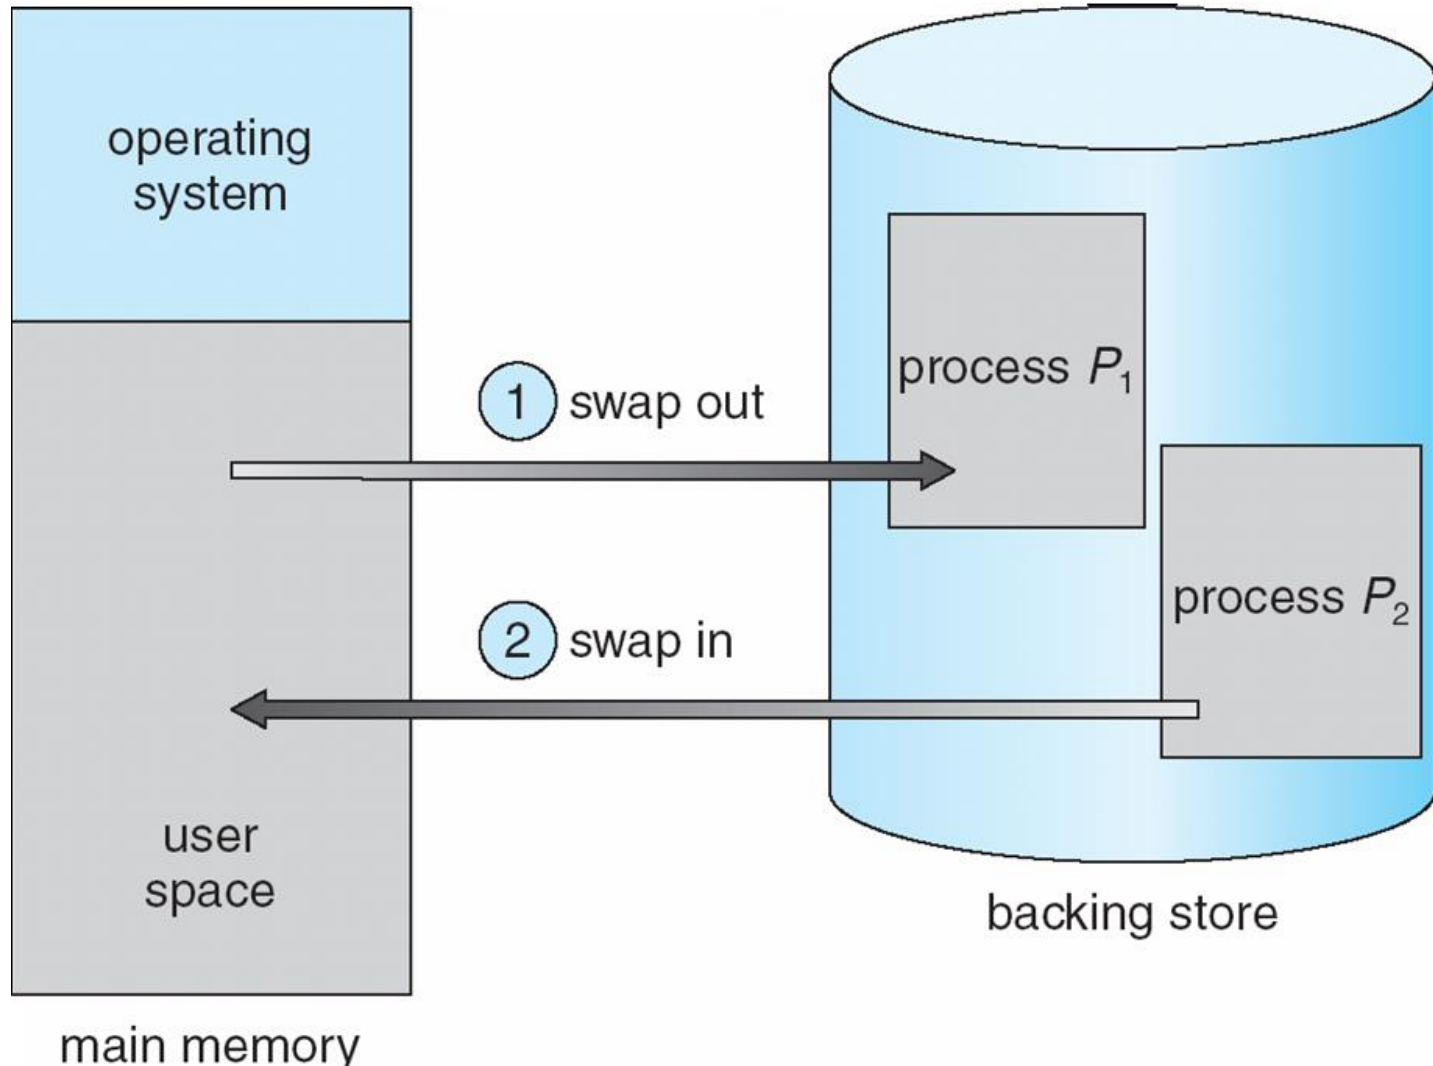
Swapping

Major part of swap time is **transfer time**; total transfer time is directly proportional to the amount of memory swapped

Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)

System maintains a **ready queue of ready-to-run processes** which have memory images on disk

Schematic View of Swapping



Contiguous Allocation

Main memory usually divides into two partitions:

Resident operating system, usually held in low memory with interrupt vector

User processes then held in high memory

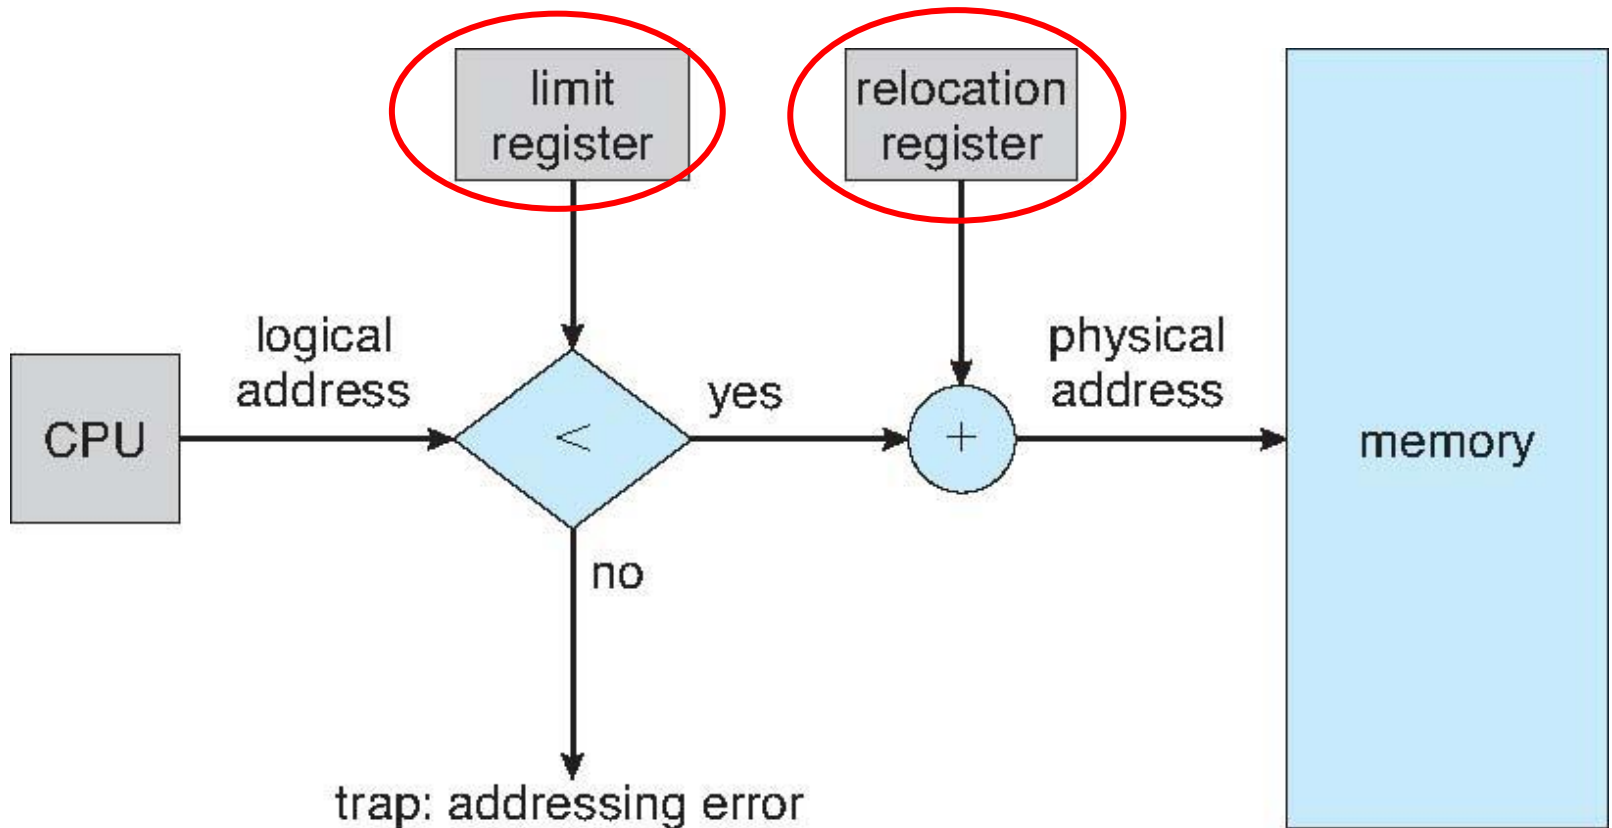
Relocation registers used to protect user processes from each other, and from changing operating-system code and data

Base register contains value of smallest physical address

Limit register contains range of logical addresses – each logical address must be less than the limit register

MMU maps logical address *dynamically*

Hardware Support for Relocation and Limit Registers



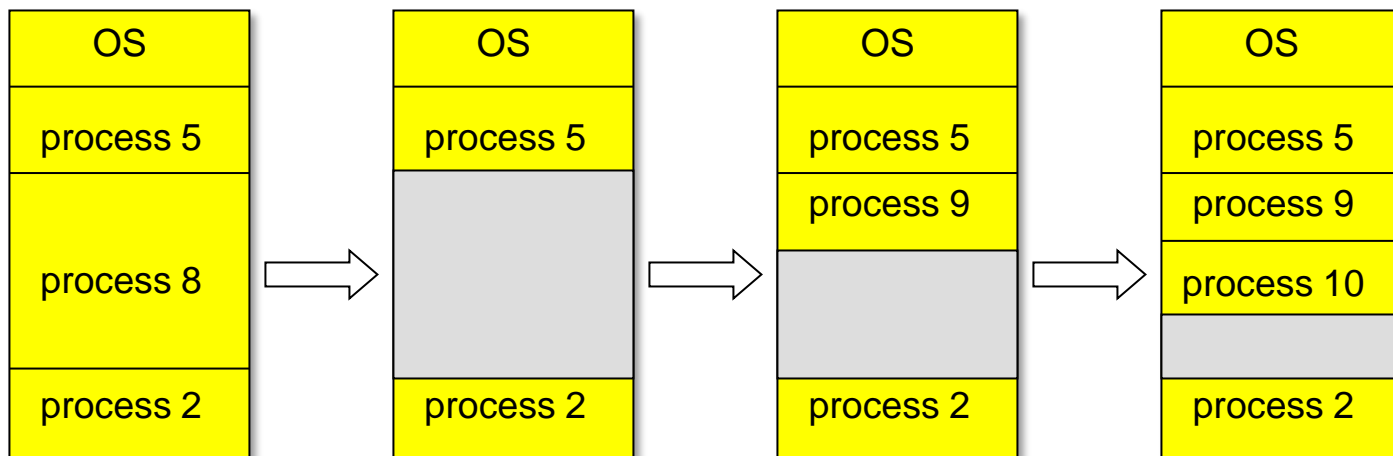
Contiguous Allocation (Cont)

Multiple-partition allocation

Hole – block of available memory; holes of various size are scattered throughout memory

When a process arrives, it is allocated memory from a hole large enough to accommodate it

Operating system maintains information about:
a) allocated partitions b) free partitions (hole)



Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes ?

First-fit: Allocate the **first hole** that is big enough

Best-fit: Allocate the **smallest hole** that is big enough;
must search entire list, unless ordered by size

Produces the smallest leftover hole

Worst-fit: Allocate the **largest hole**; must also search
entire list

Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed
and storage utilization

Fragmentation

External Fragmentation – total memory space exists to satisfy a request, but it is not contiguous

Internal Fragmentation – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used

Reduce external fragmentation by **compaction**

Shuffle memory contents to place all free memory together in one large block

Compaction is possible *only* if relocation is dynamic, and is done at execution time

I/O problem

- ▶ Latch job in memory while it is involved in I/O
- ▶ Do I/O only into OS buffers

Paging

Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available

Paging avoids external fragmentation and the needs for compaction.

Divide **physical memory** into fixed-sized blocks called **frames** (size is power of 2, between 512 - 8,192 bytes)

Divide **logical memory** into blocks of same size called **pages**

Keep track of all free frames

To run a program of size **n pages**, need to find **n free frames** and load program

Set up a **page table to translate logical to physical addresses**


Internal fragmentation

Address Translation Scheme

Address generated by CPU is divided into:

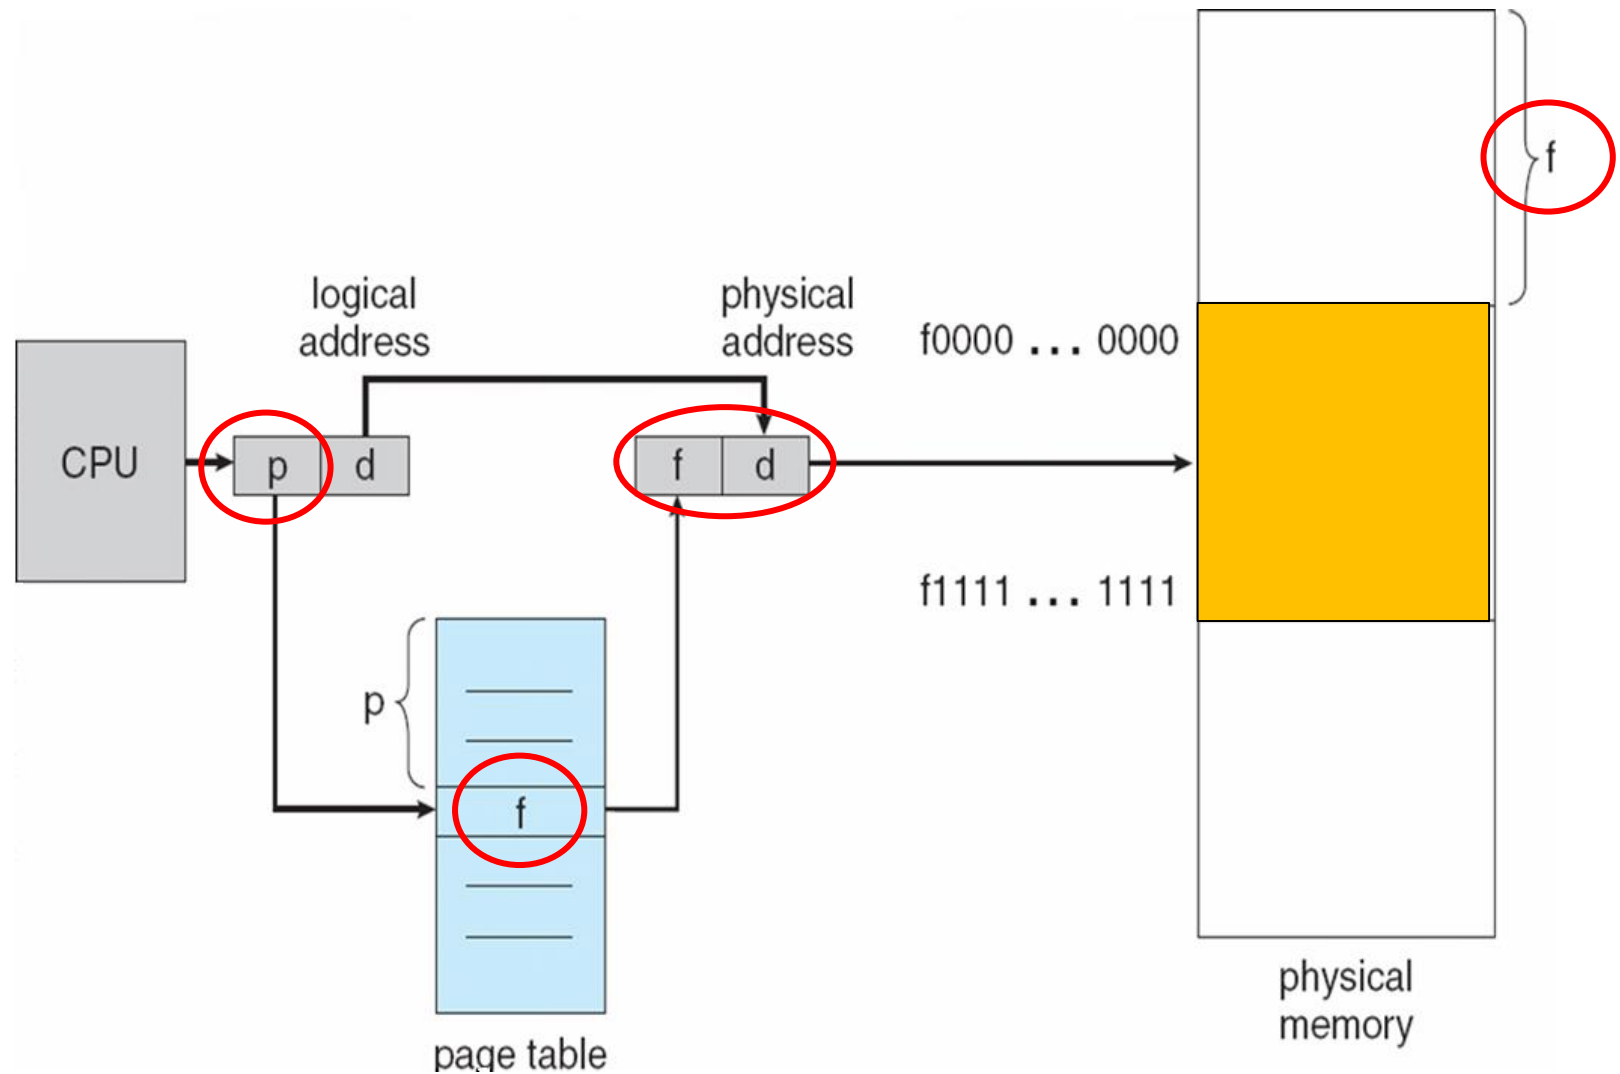
Page number (p) – used as an index into a **page table** which contains base address of each page in physical memory

Page offset (d) – combined with base address to define the physical memory address that is sent to the memory unit

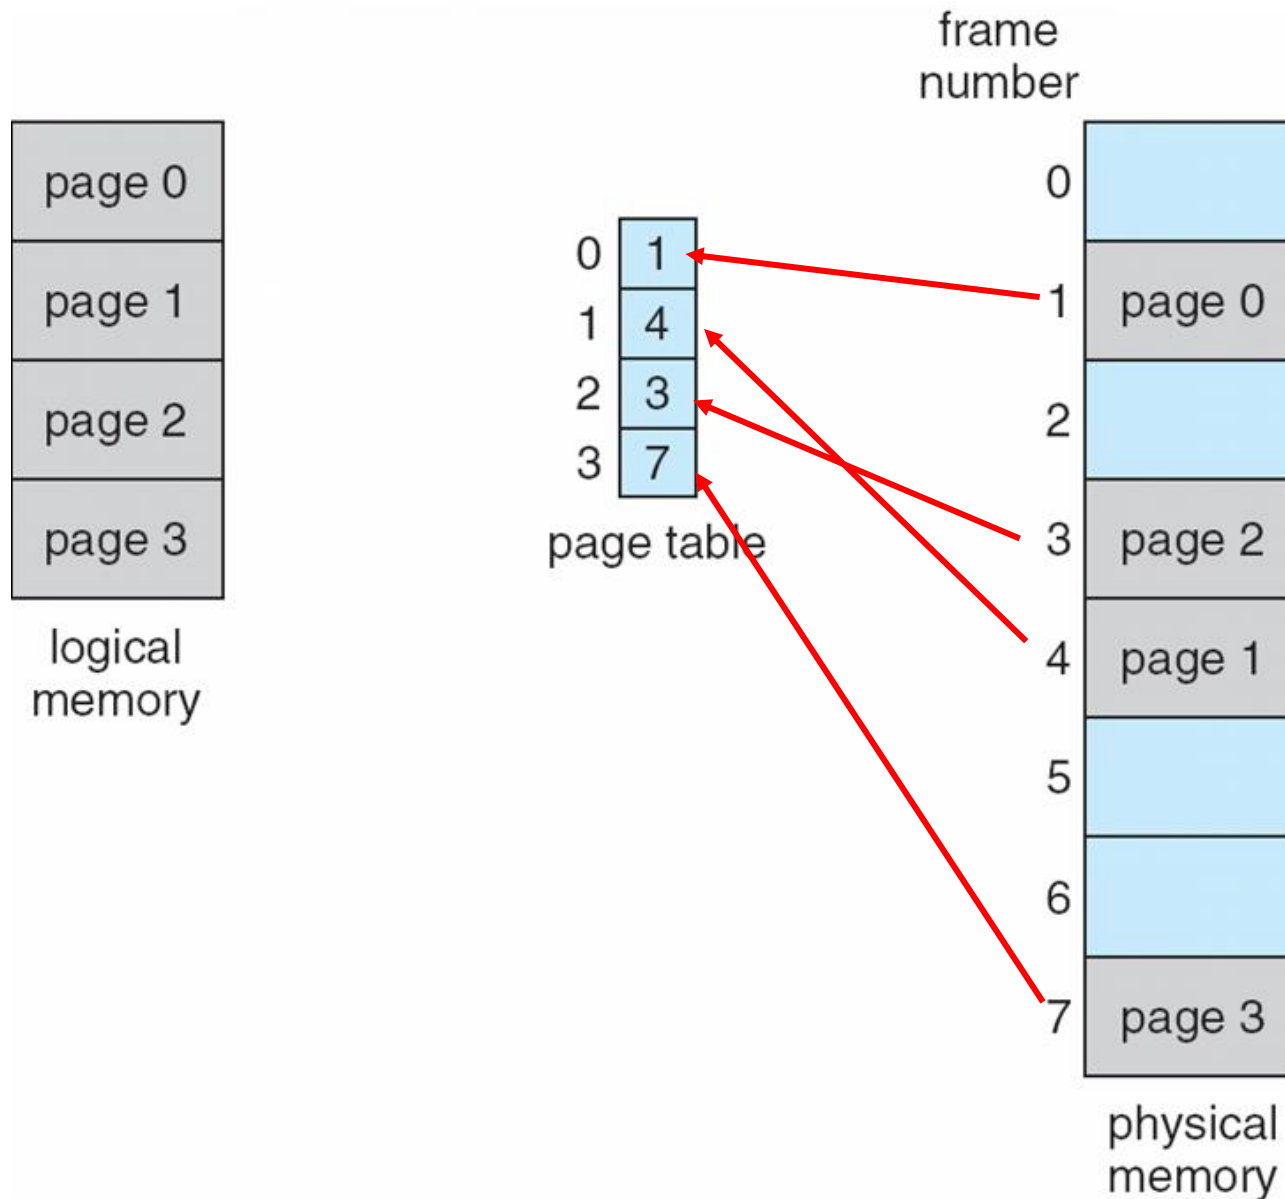


For given **logical address space 2^m** and **page size 2^n**


Paging Hardware



Paging Model of Logical and Physical Memory



Paging Example



32-byte memory and 4-byte pages

Free Frames



Before allocation

After allocation

Implementation of Page Table

Page table is kept in main memory

Page-table base register (PTBR) points to the page table

Page-table length register (PRLR) indicates size of the page table

In this scheme every data/instruction access requires **two memory accesses**. One for the page table and one for the data/instruction.

The two memory access problem can be solved by the use of a special **fast-lookup hardware cache** called **associative memory** or **translation look-aside buffers (TLBs)**

Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – **uniquely identifies each process** to provide address-space protection for that process

Associative Memory

Associative memory – provides **parallel search**

$p = 6$

Page #	Frame #
3	8
4	10
6	5
10	7

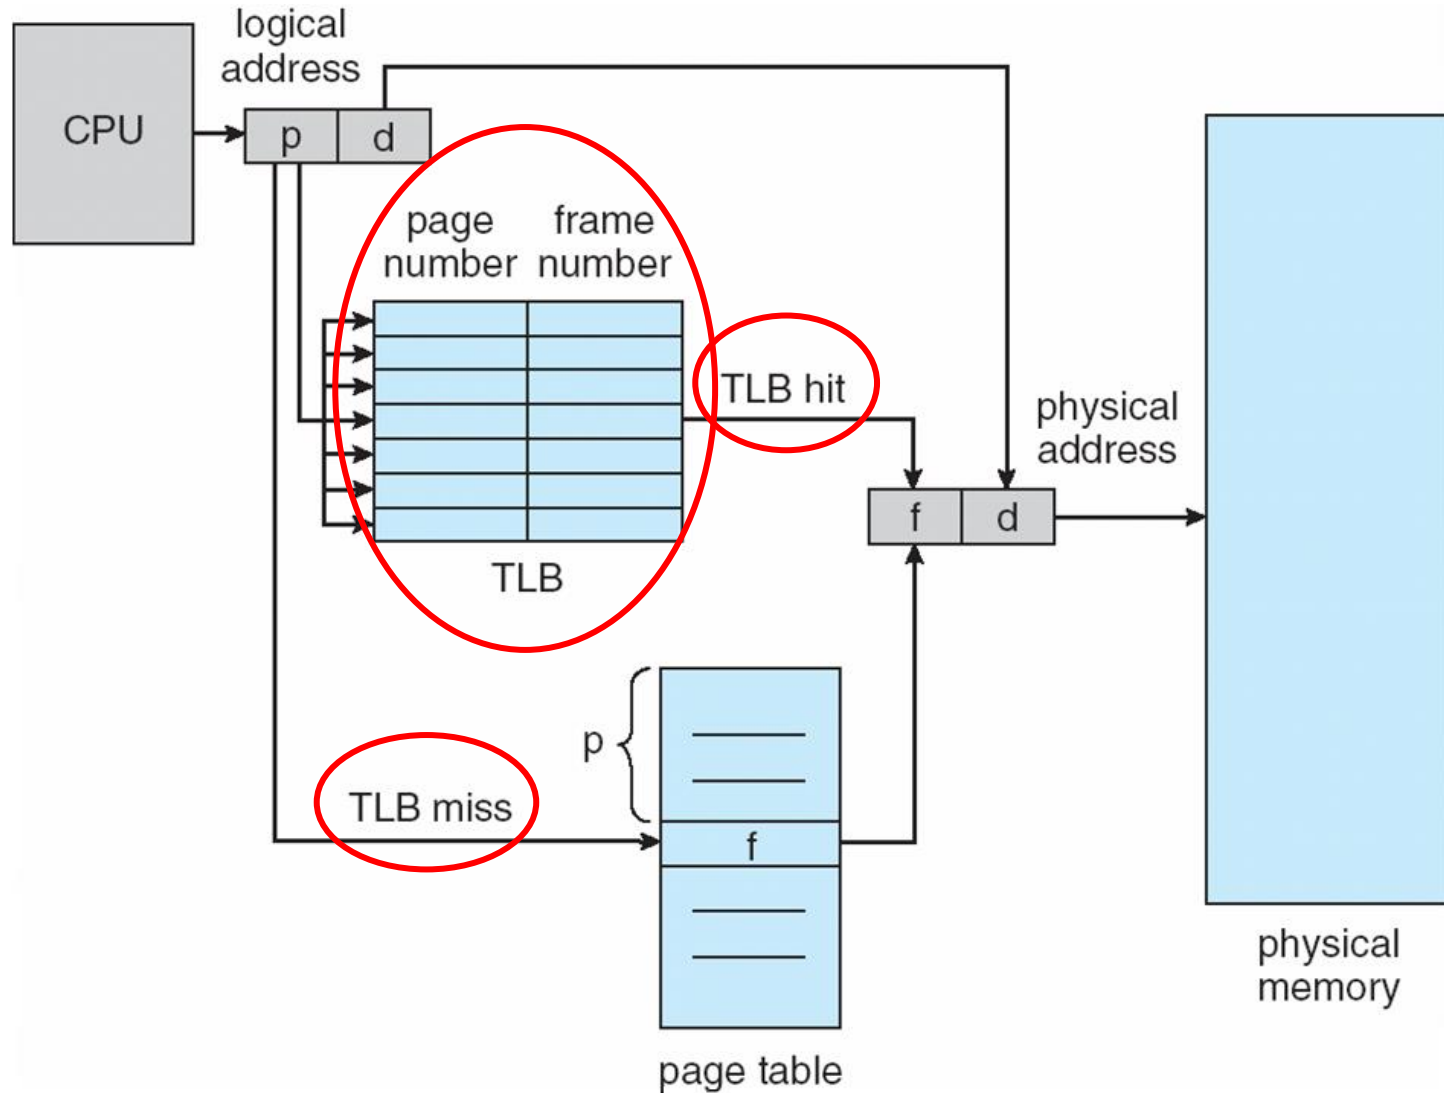
$F\# = 5$

■ Address translation (p, d)

If p is in associative register, get frame # out

Otherwise get frame # from page table in memory

Paging Hardware With TLB



Effective Access Time

Associative Lookup = ϵ time unit

Assume memory cycle time is 1 microsecond

Hit ratio – percentage of times that a page number is found in the associative registers;
ratio related to number of associative registers

Hit ratio = α

Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} &= (1 + \epsilon) \alpha + (2 + \epsilon)(1 - \alpha) \\ &= 2 + \epsilon - \alpha \end{aligned}$$

Memory Protection

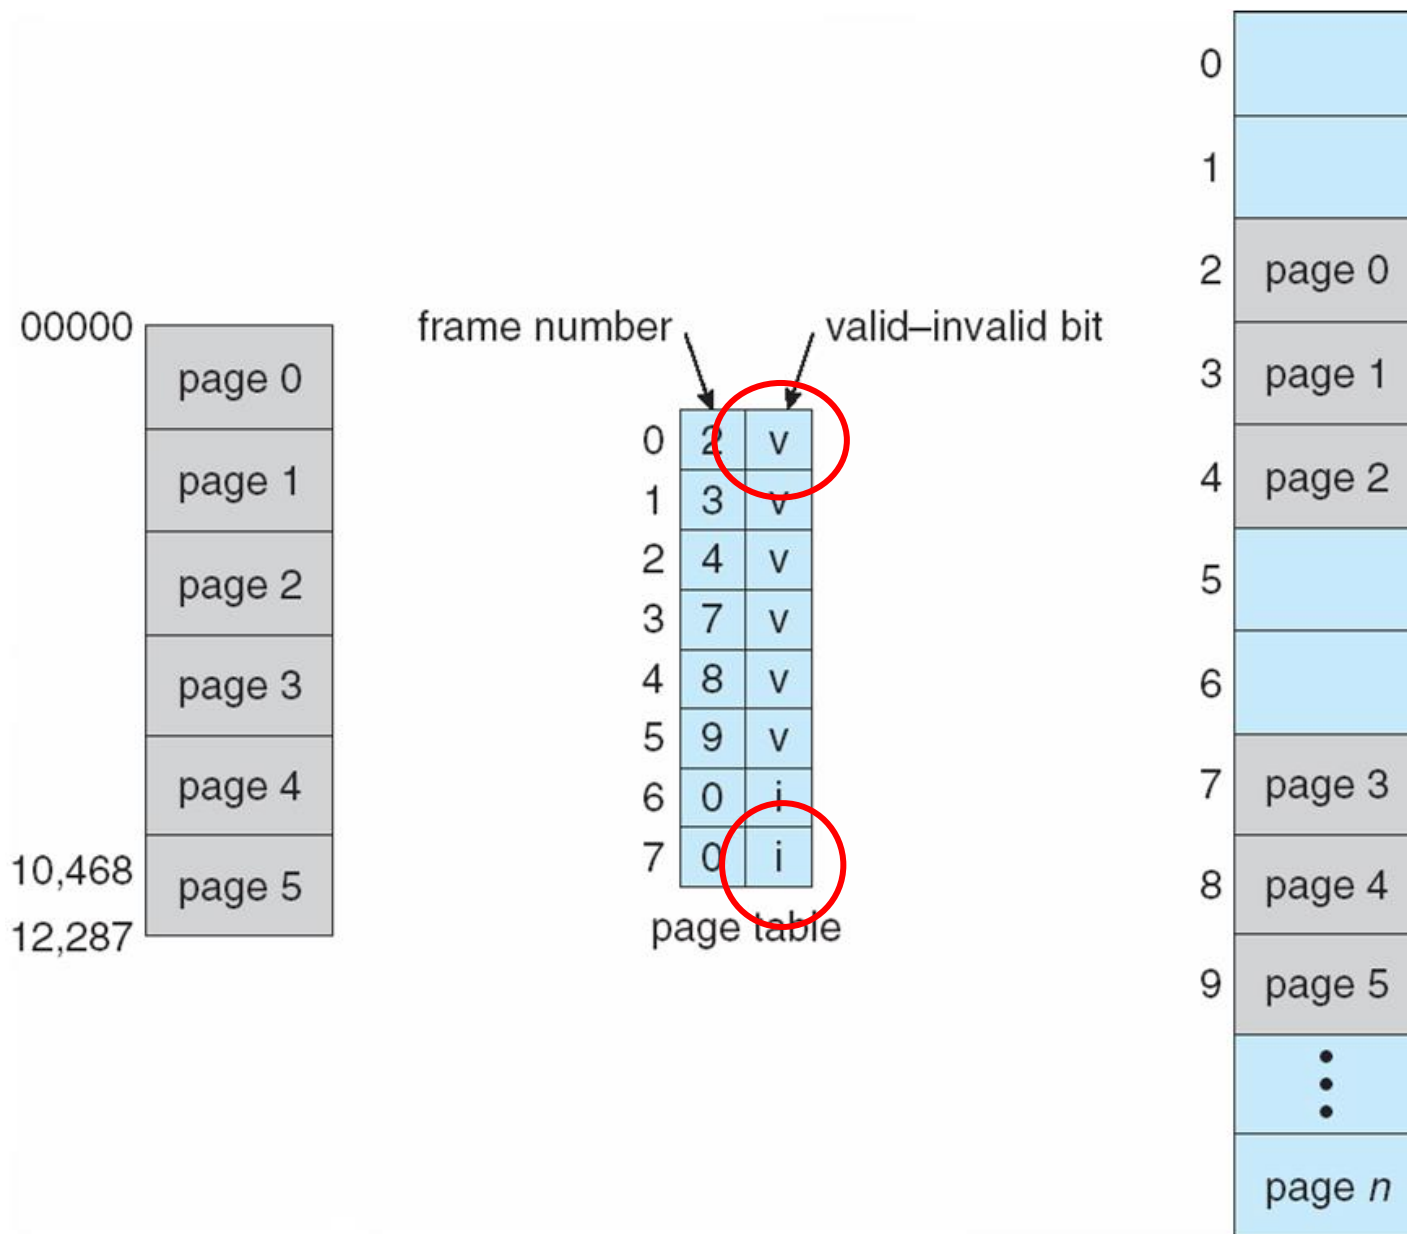
Memory protection implemented by associating **protection bit** with each frame

Valid-invalid bit attached to each entry in the page table:

“valid” indicates that the associated page is **in the process’ logical address space**, and is thus a legal page

“invalid” indicates that the page is not in the process’ logical address space

Valid (v) or Invalid (i) bit in a Page Table



Shared Pages

Shared code

One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).

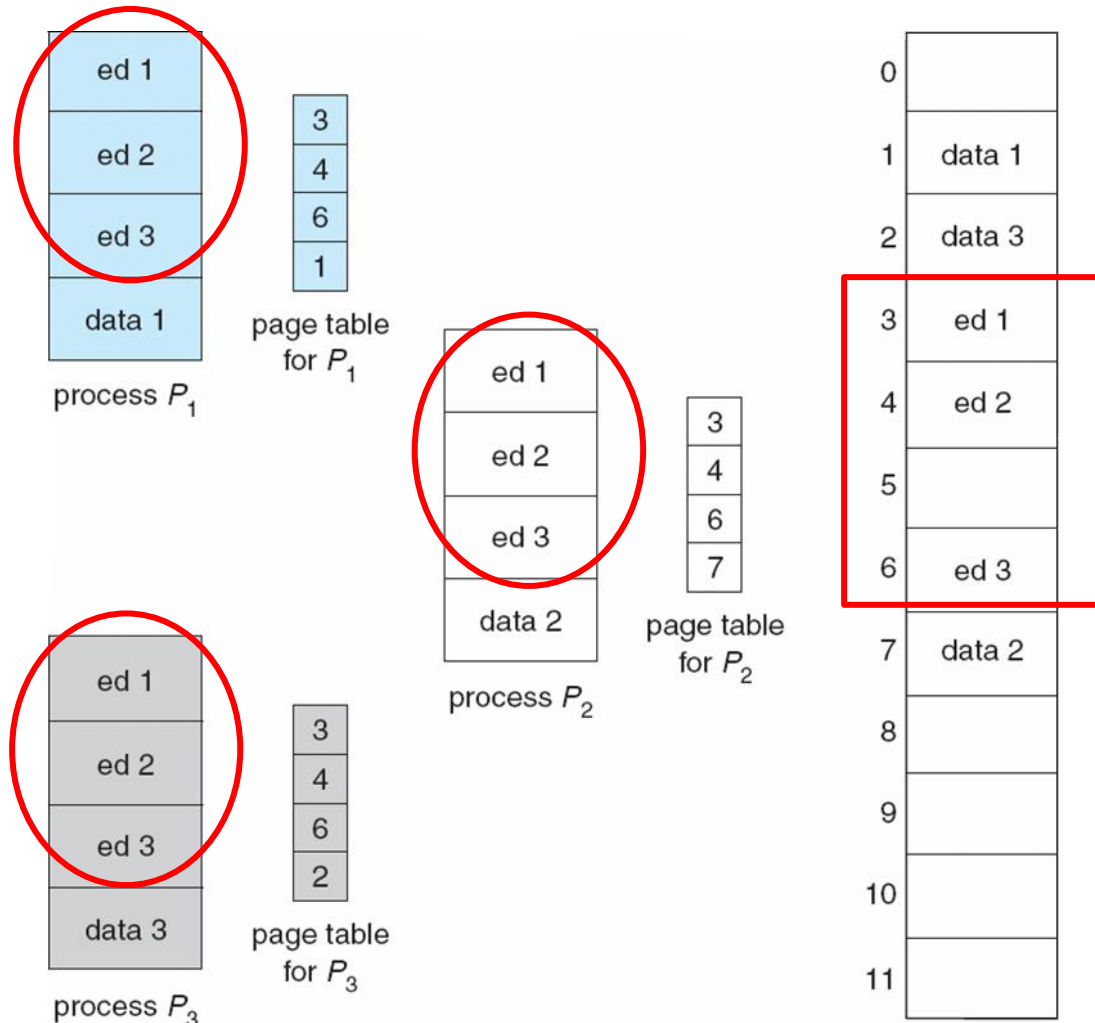
Reentrant code is non-self-modifying code: it never changes during execution.

Private code and data

Each process keeps a separate copy of the code and data

Some operating systems implement shared memory using shared pages.

Shared Pages Example



Structure of the Page Table

Hierarchical Paging

Hashed Page Tables

Inverted Page Tables

Hierarchical Page Tables


Modern computer systems support a large logical address space 2^{32} to 2^{64} . **The page table itself becomes excessively large.**

For 32-bit logical address space, and page size of 4K, then a page table consists of 1 million entries ($2^{32} / 2^{12} = 2^{20} = 1 \text{ million}$).

Break up the logical address space into multiple page tables

A simple technique is a **two-level page table**

Two-Level Page-Table Scheme



Two-Level Paging Example

A logical address (on 32-bit machine with 1K page size) is divided into:

a **page number** consisting of 22 bits


a **page offset** consisting of 10 bits

Since the page table is paged, the page number is further divided into:

a 12-bit page number

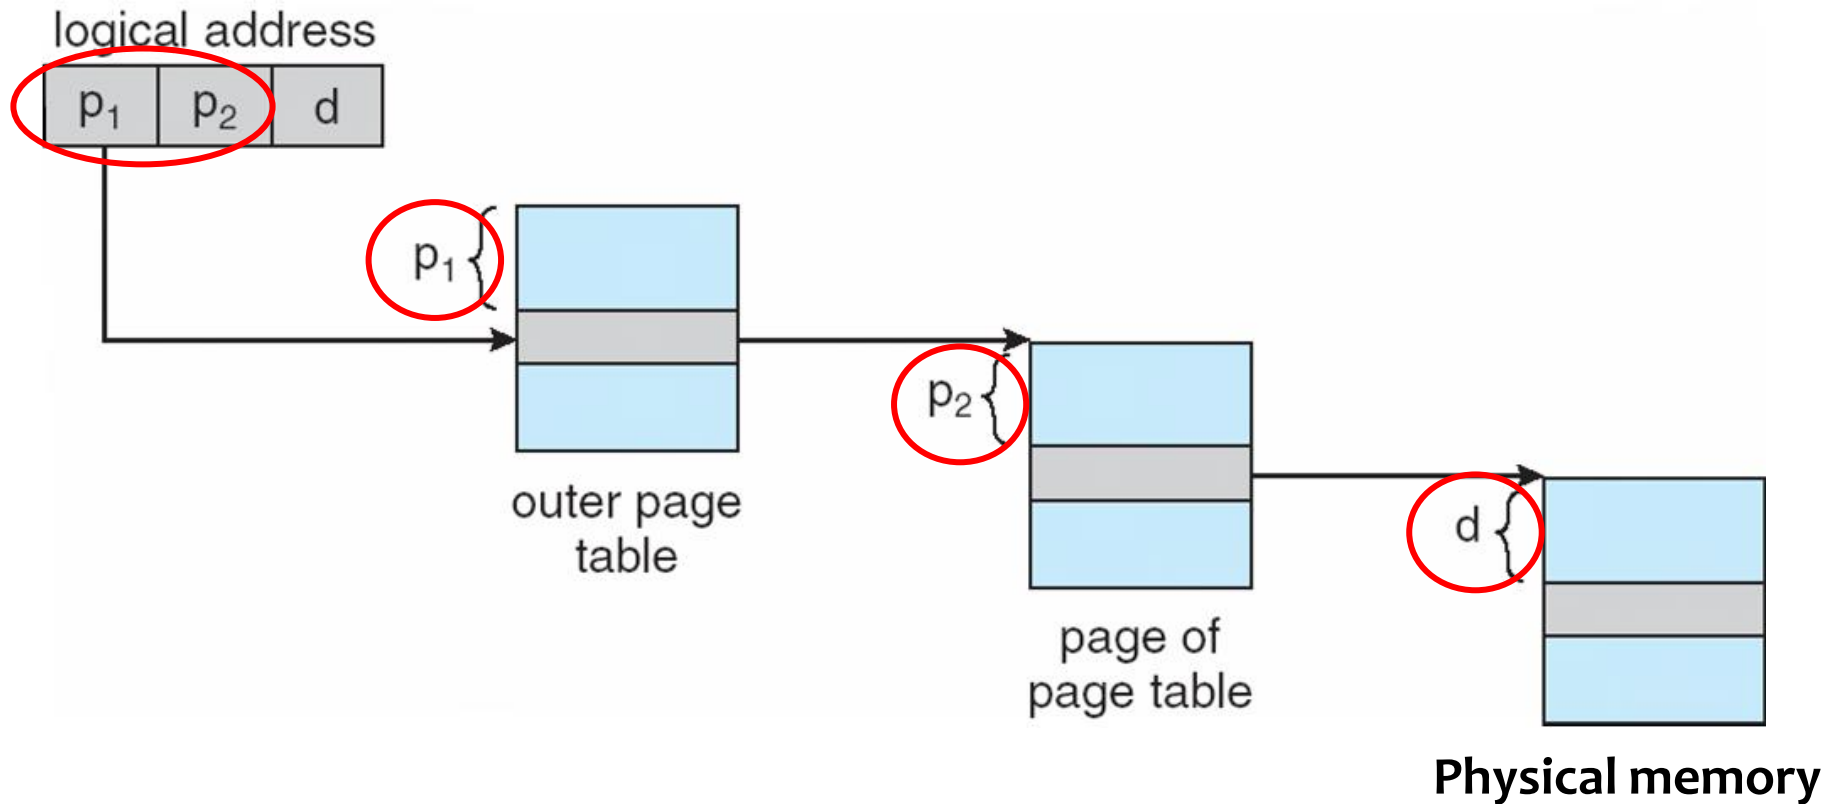
a 10-bit page offset

Thus, a logical address is as





where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

Address-Translation Scheme



Three-level Paging Scheme



64-bit machine with 4K page

Hashed Page Tables

Common in address spaces > 32 bits


The **virtual page number is hashed into** a page table

This page table contains **a chain of elements hashing to the same location**

Virtual page numbers are compared in this chain searching for a match

If a match is found, the corresponding physical frame is extracted

Hashed Page Table



Inverted Page Table

One entry for each real page of memory


The page table is shared by all processes

Entry consists of the **virtual address of the page stored in that real memory location**, with information about the **process** that owns that page

Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs

Use hash table to limit the search to one — or at most a few — page-table entries

Inverted Page Table Architecture



The search can be done sequentially, or
by hash function, or
by associative memory

Segmentation

Memory-management scheme that supports **user view of memory**

A program is **a collection of segments**

A segment is a logical unit such as:

- main program

- procedure

- function

- method

- object

- local variables, global variables


- common block

- stack


- symbol table

- arrays


User's View of a Program



Logical View of Segmentation



user space



physical memory space

Segmentation Architecture

Logical address consists of a two-tuple:

<segment-number, offset>,

Segment table – maps two-dimensional physical addresses; each table entry has:

base – contains the starting physical address where the segments reside in memory

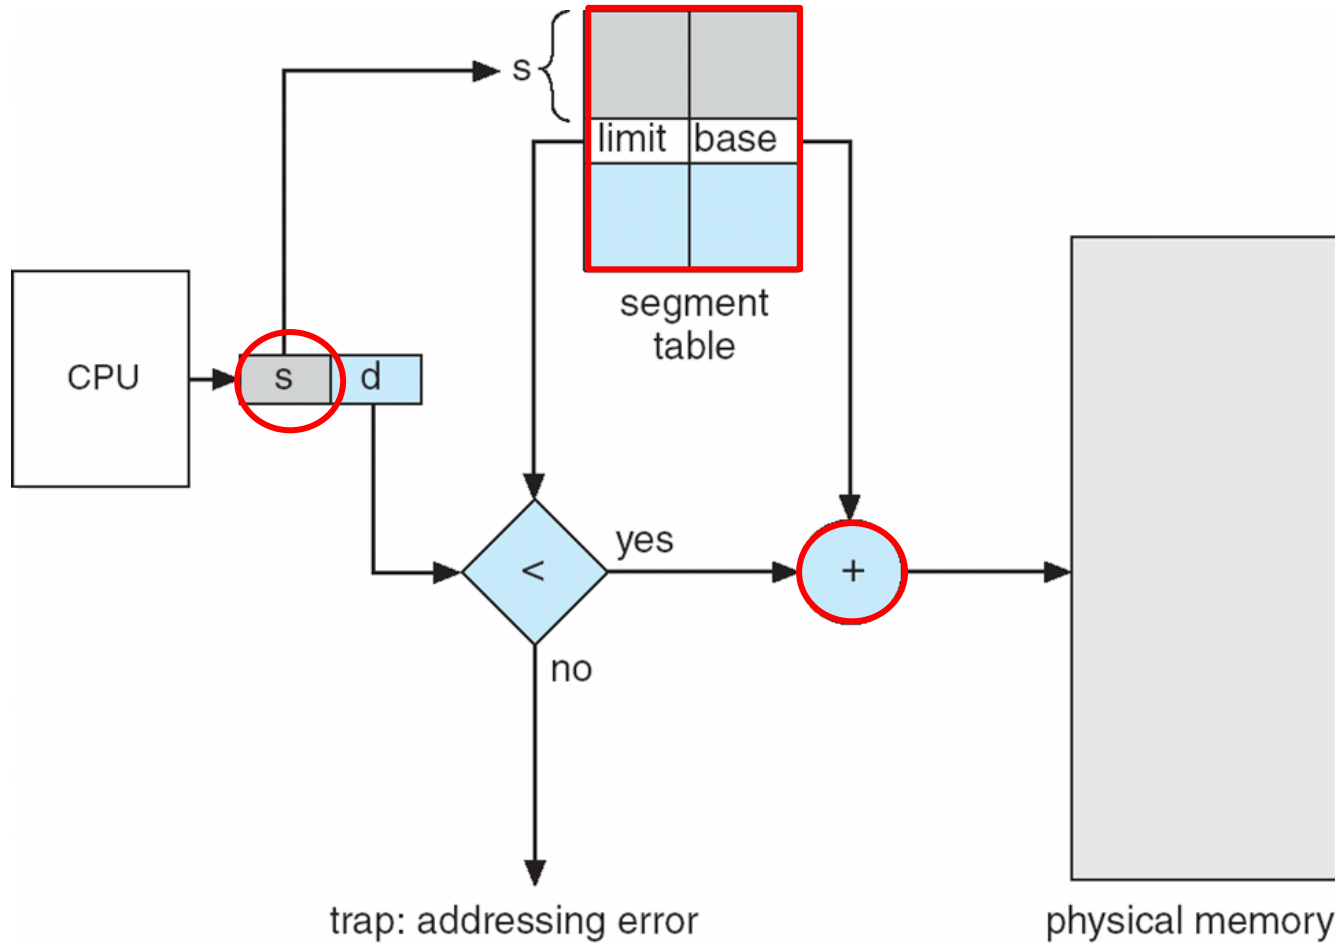
limit – specifies the length of the segment

Segment-table base register (STBR) points to the segment table's location in memory

Segment-table length register (STLR) indicates **number of segments** used by a program;

segment number **s** is legal if **s < STLR**

Segmentation Hardware



Segmentation Architecture (Cont.)

Protection

With each entry in segment table associate:

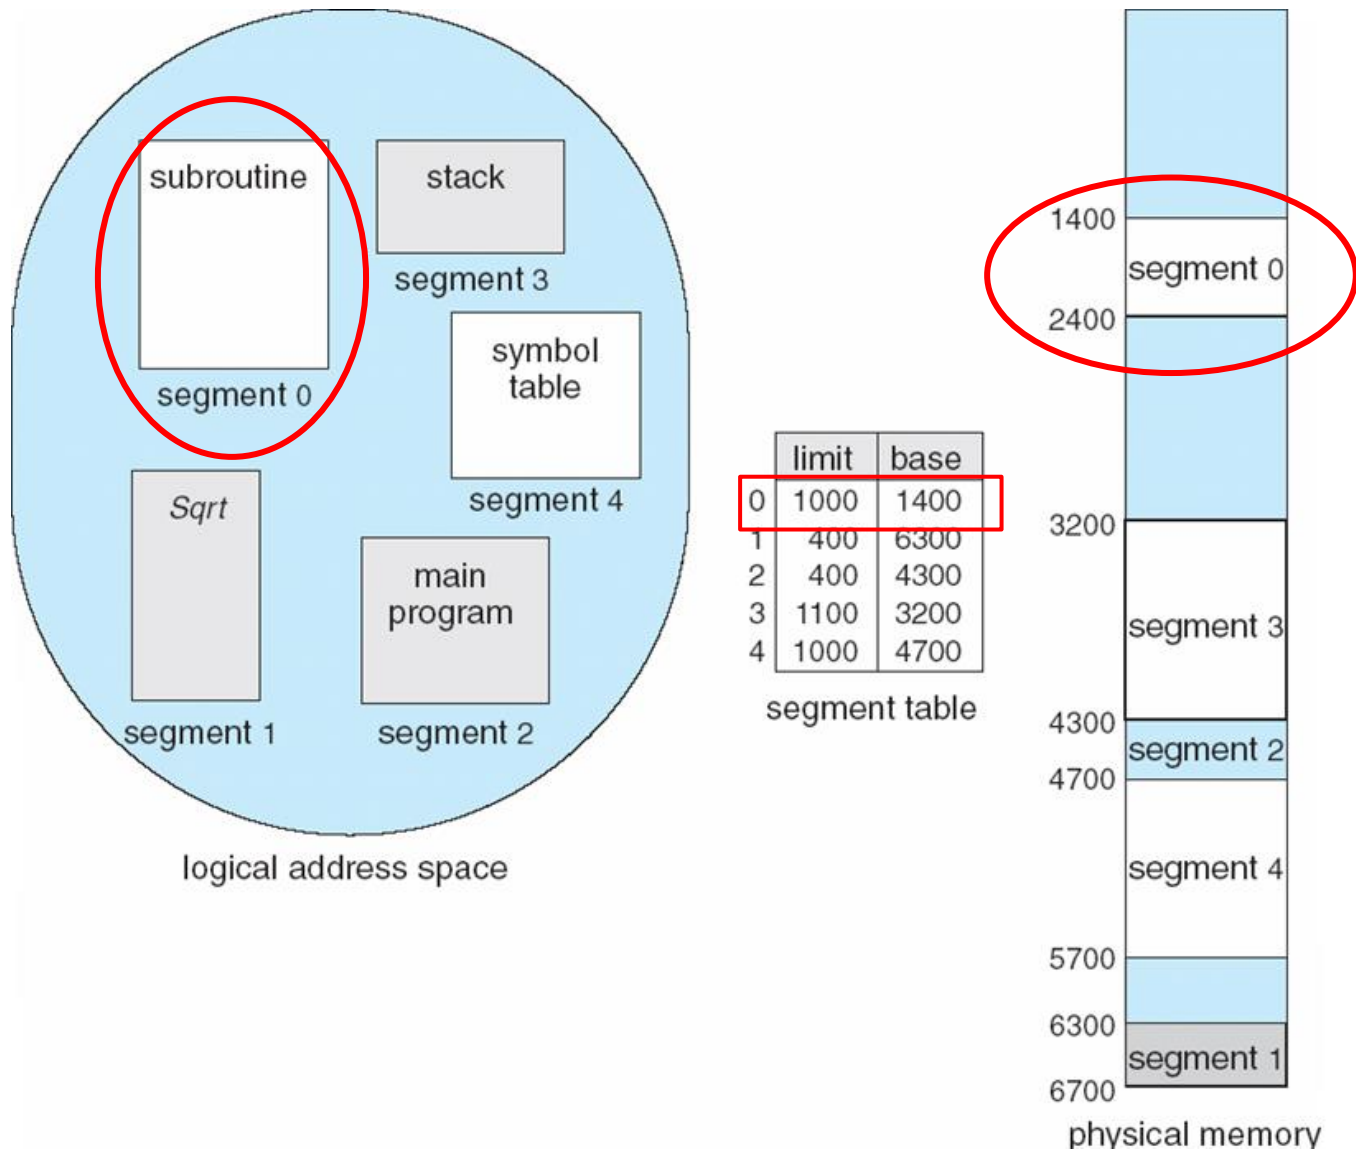
- ▶ validation bit = 0 \Rightarrow illegal segment
- ▶ read/write/execute privileges

Protection bits associated with segments; code sharing occurs at segment level

Since segments vary in length, **memory allocation is a dynamic storage-allocation problem**

A segmentation example is shown in the following diagram

Example of Segmentation



Example: The Intel Pentium

Supports both **segmentation** and **segmentation with paging**


CPU generates **logical address**

Given to segmentation unit

- ▶ Which produces **linear addresses**

Linear address given to paging unit

- ▶ Which generates physical address in main memory
- ▶ Paging units form equivalent of MMU



Pentium Segmentation

A segment is allowed to be **as large as 4GB (32 bits)**, and the maximum number of segments per process is **16K**.

The logical address space is divided into **two partitions**:

The first partition up to **8K segments** that are private to that process.

The 2nd partition up to **8K segments** that are shared among all processes.

Local Descriptor Table (LDT): Information about the 1st partition

Global Descriptor Table (GDT): Information about the 2nd partition.

Each entry of LDT and GDT is an **8-byte descriptor** of a particular segment.

Pentium Segmentation

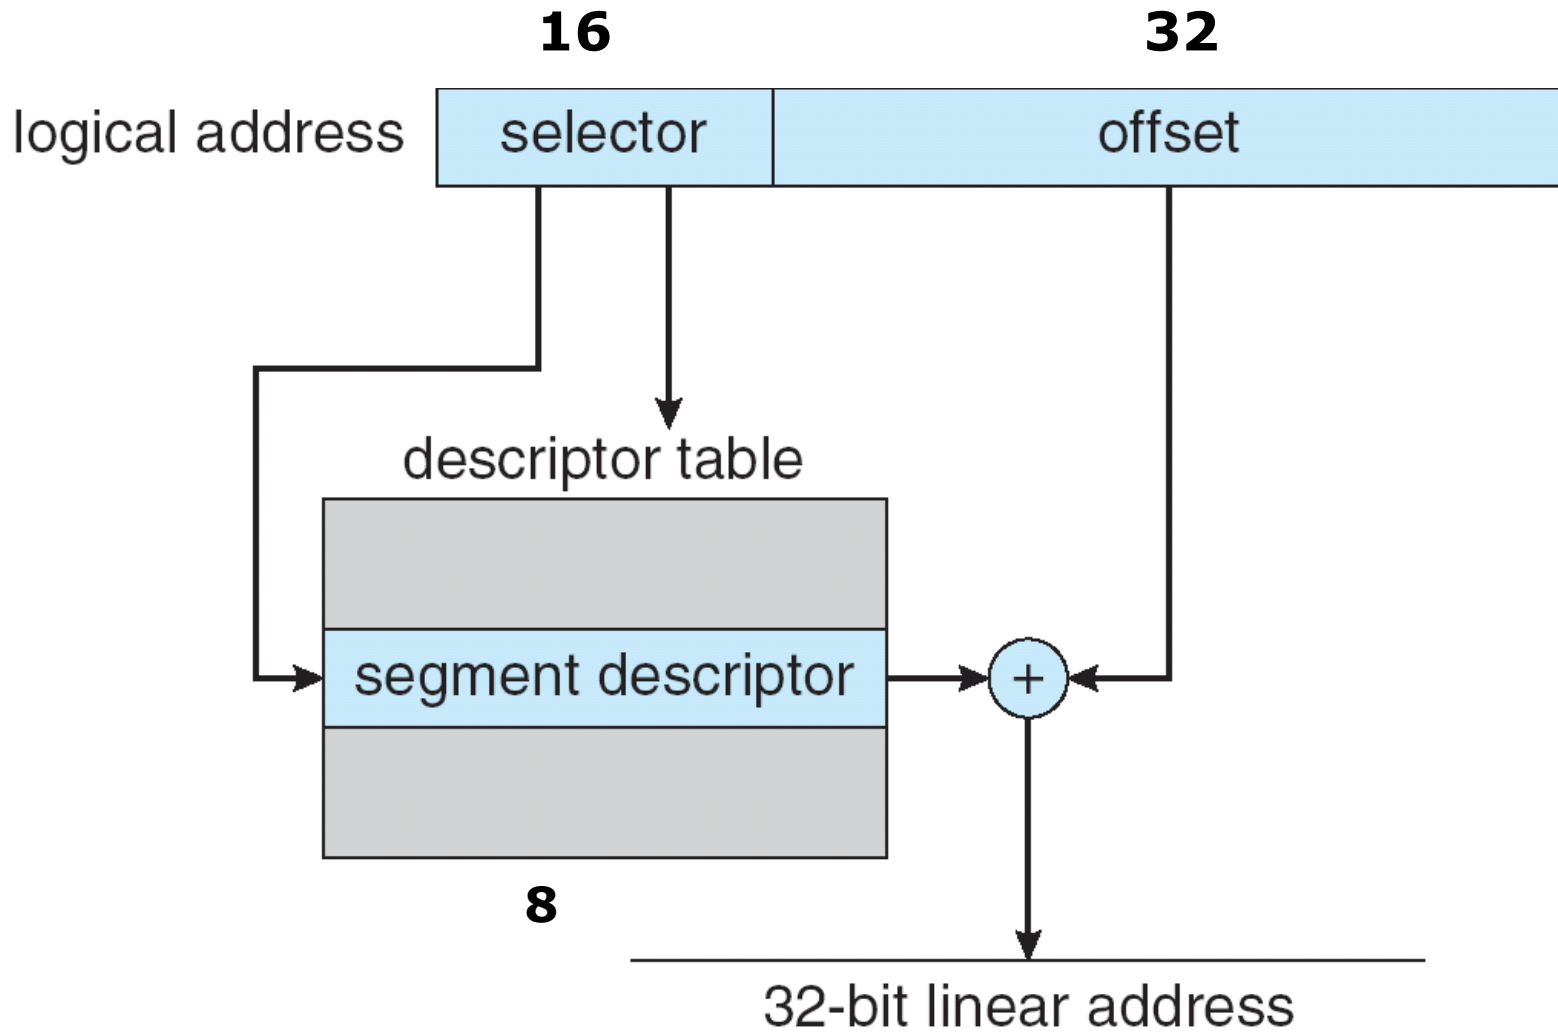
The logical address is a pair (**selector, offset**), where the **selector** is a 16-bit number and offset is a 32-bit number:



The machine has **six segment registers**, allowing six segments to be addressed at any one time by a process.

The **linear address is 32 bits long** and is formed as follows.


Intel Pentium Segmentation



Pentium Paging

A page is allowed to be as 4kB or 4MB.

For 4-KB pages, a two level paging scheme is used (32-bit linear address)

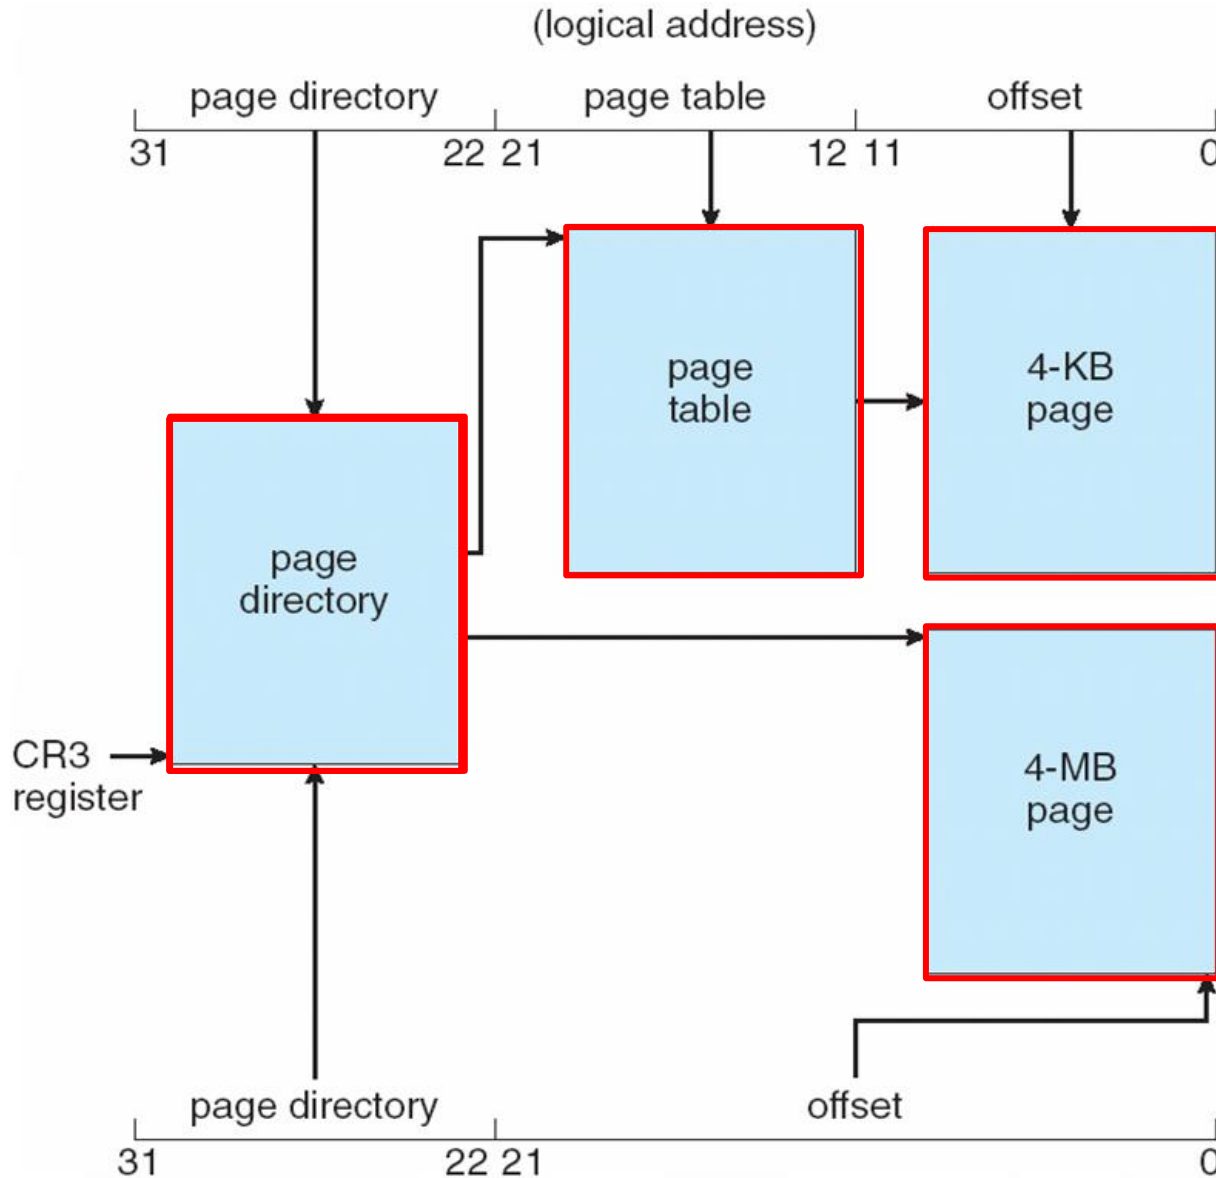


The address-translation scheme is shown as follows.

Page directory

Page size flag

Pentium Paging Architecture



Linux on Pentium Systems

Linux does not rely on segmentations and used it minimally.

On the Pentium, **Linux uses only six segments:**

- A segment for kernel code

- A segment for kernel data

- A segment for user code

- A segment for user data

- A task-state segment (TSS)

- A default LDT segment

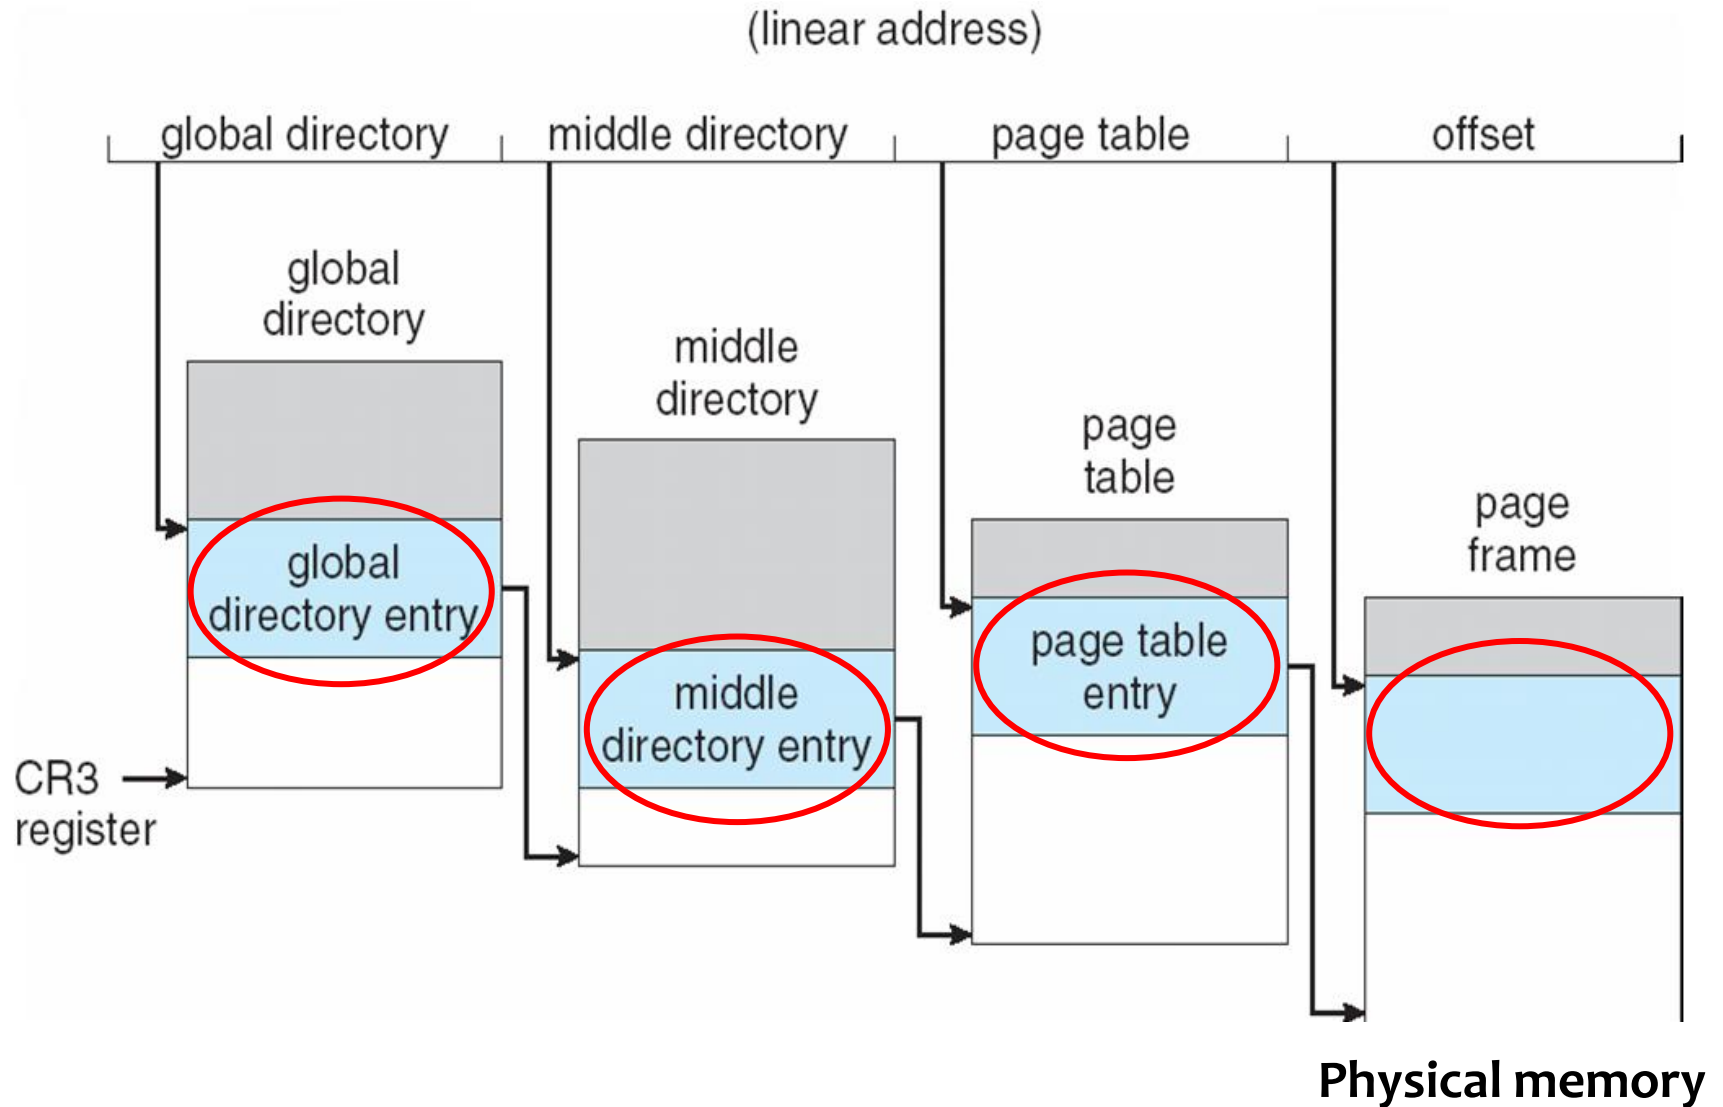
Linear Address in Linux

- The linear address in Linux is broken **into four parts**:

global directory	middle directory	page table	offset
---------------------	---------------------	---------------	--------

- Each task in Linux has its own set of page tables and the **CR3 register points to the global directory** for the task currently executing.
- During a **context switch**, the value of CR3 register is restored in the **TSS segments of the tasks** involved in the context switch.

Three-level Paging in Linux



End of Chapter 8

